


llama.cpp

Model Files (GGUF)



Core/Runtime

- Arch definitions
- Model loading
- Tokenizer
- Graph construction
- Manage KV-cache
- Backend dispatch

Backends

CUDA

Metal

WebGPU

Targets

NVIDIA

Apple

